



Федеральное государственное бюджетное образовательное учреждение  
высшего профессионального образования  
«Московский государственный технический университет имени Н.Э.  
Баумана»  
(МГТУ им. Н.Э. Баумана)

**Лабораторная работа 4**  
**по курсу «Технологии машинного обучения»**

Выполнил  
студент группы ИУ5-64  
XXX

Москва, 2021

## **Цель работы**

изучение линейных моделей, SVM и деревьев решений.

## **Задание**

1. Выберите набор данных (датасет) для решения задачи классификации или регрессии.
2. В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.
3. С использованием метода `train_test_split` разделите выборку на обучающую и тестовую.
4. Обучите следующие модели:
  - одну из линейных моделей;
  - SVM;
  - дерево решений.
5. Оцените качество моделей с помощью двух подходящих для задачи метрик. Сравните качество полученных моделей.

## **Дополнительные задания**

1. Проведите эксперименты с важностью признаков в дереве решений.
2. Визуализируйте дерево решений.

## **Ход работы**

```
In [1]: # Загрузка данных
import pandas as pd
data = pd.read_csv("../2/melbourne_housing.csv")
```

```
In [2]: display(data.dtypes), display(data.head()), display(data.isnull().sum());
```

Suburb object  
Address object  
Rooms int64  
Type object  
Price float64  
Method object  
SellerG object  
Date object  
Distance float64  
Postcode float64  
Bedroom2 float64  
Bathroom float64  
Car float64  
Landsize float64  
BuildingArea float64  
YearBuilt float64  
CouncilArea object  
Latitude float64  
Longitude float64  
Regionname object  
Propertycount float64  
dtype: object

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	Distance	Postcode	...	Bathroom	Car	Landsize	BuildingArea
0	Abbotsford	68 Studley St	2	h	NaN	SS	Jellis	3/09/2016	2.5	3067.0	...	1.0	1.0	126.0	NaN
1	Abbotsford	85 Turner St	2	h	1480000.0	S	Biggin	3/12/2016	2.5	3067.0	...	1.0	1.0	202.0	NaN
2	Abbotsford	25 Bloomburg St	2	h	1035000.0	S	Biggin	4/02/2016	2.5	3067.0	...	1.0	0.0	156.0	79.0
3	Abbotsford	18/659 Victoria St	3	u	NaN	VB	Rounds	4/02/2016	2.5	3067.0	...	2.0	1.0	0.0	NaN
4	Abbotsford	5 Charles St	3	h	1465000.0	SP	Biggin	4/03/2017	2.5	3067.0	...	2.0	0.0	134.0	150.0

5 rows × 21 columns

Suburb	0
Address	0
Rooms	0
Type	0
Price	7610
Method	0
SellerG	0
Date	0
Distance	1
Postcode	1
Bedroom2	8217
Bathroom	8226
Car	8728
Landsize	11810
BuildingArea	21115
YearBuilt	19306
CouncilArea	3
Latitude	7976
Longitude	7976
Regionname	3
Propertycount	3
dtype:	int64

```
In [3]: import numpy as np
from sklearn.preprocessing import LabelEncoder

columns_and_types = {
    "Rooms": np.int64,
    "Type": None,
    "Price": np.int64,
    "Distance": np.float64,
```

```

"Postcode": np.int64,
"Bedroom2": np.int64,
"Bathroom": np.int64,
"Car": np.int64,
"Landsize": np.float64,
"BuildingArea": np.float64,
"YearBuilt": np.int64,
"Lattitude": np.float64,
"Longtitude": np.float64,
"Propertycount": np.int64,
}

data = data[list(columns_and_types.keys())]
data.dropna(axis=0, how='any', inplace=True)
data = data.astype({k: v for k,v in columns_and_types.items() if v is not None})

type_encoder = LabelEncoder()
data["Type"] = type_encoder.fit_transform(data["Type"])

```

In [4]: `display(data.shape), display(data.dtypes)`

(8887, 14)

```

Rooms          int64
Type            int64
Price           int64
Distance        float64
Postcode        int64
Bedroom2        int64
Bathroom        int64
Car             int64
Landsize        float64
BuildingArea    float64
YearBuilt       int64
Lattitude       float64
Longtitude      float64
Propertycount   int64
dtype: object

```

Out[4]: (None, None)

In [5]: `data.head()`

Out[5]:

	Rooms	Type	Price	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	BuildingArea	YearBuilt	Lattitude	Longtitude	Propertycount
2	2	0	1035000	2.5	3067	2	1	0	156.0	79.0	1900	-37.8079	144.9934	4
4	3	0	1465000	2.5	3067	3	2	0	134.0	150.0	1900	-37.8093	144.9944	4
6	4	0	1600000	2.5	3067	3	1	2	120.0	142.0	2014	-37.8072	144.9941	4
11	3	0	1876000	2.5	3067	4	2	0	245.0	210.0	1910	-37.8024	144.9993	4
14	2	0	1636000	2.5	3067	2	1	2	256.0	107.0	1890	-37.8060	144.9954	4

In [6]:

```

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
for col in data.columns:
    if col != "Price":
        data[col] = scaler.fit_transform(data[[col]])

```

In [7]: `data.head()`

Out[7]:

	Rooms	Type	Price	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	BuildingArea	YearBuilt	Lattitude	Longtitude	Propertycount
2	-1.140264	-0.556327	1035000	-1.27695	-0.396621	-1.115905	-0.895892	-1.734910	-0.346267	-0.799693	-1.775256	-0.037540	0.0168	4
4	-0.102631	-0.556327	1465000	-1.27695	-0.396621	-0.080939	0.489973	-1.734910	-0.366997	0.007854	-1.775256	-0.053003	0.0252	4
6	0.935002	-0.556327	1600000	-1.27695	-0.396621	-0.080939	-0.895892	0.315512	-0.380188	-0.083137	1.302598	-0.029809	0.0227	4
11	-0.102631	-0.556327	1876000	-1.27695	-0.396621	0.954028	0.489973	-1.734910	-0.262404	0.690288	-1.505269	0.023204	0.0664	4
14	-1.140264	-0.556327	1636000	-1.27695	-0.396621	-1.115905	-0.895892	0.315512	-0.252039	-0.481224	-2.045243	-0.016556	0.0336	4


In [8]:

```
import seaborn as sns
import matplotlib.pyplot as plt

fig, ax = plt.subplots(figsize=(15,7))
sns.heatmap(data.corr(method="pearson"), ax=ax, annot=True, fmt='.2f');
```



```
In [9]: fig, ax = plt.subplots(figsize=(10,10))
sns.scatterplot(ax=ax, x="BuildingArea", y="Price", data=data);
```



## Линейная регрессия

```
In [10]: from sklearn.model_selection import train_test_split
```

```
data_X = data.loc[:, [x for x in data.columns if x != "Price"]]
data_Y = data.loc[:, 'Price']
data_X_train, data_X_test, data_y_train, data_y_test = train_test_split(
    data_X,
    data_Y,
    test_size=0.2,
    random_state=1
)
```

```
In [11]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_percentage_error

# Обучим линейную регрессию и сравним коэффициенты с рассчитанными ранее
reg1 = LinearRegression()
reg1.fit(data_X_train, data_y_train)
y_pred = reg1.predict(data_X_test)
r2_score(data_y_test, y_pred), mean_absolute_percentage_error(data_y_test, y_pred)
```

```
Out[11]: (0.6175514371440386, 0.26894243290410874)
```

## SVR

```
In [12]: from sklearn.svm import SVR

svr = SVR(max_iter=10000, kernel="rbf", C=1e6)
svr.fit(data_X_train, data_y_train);
```

```
/usr/lib/python3.9/site-packages/sklearn/svm/_base.py:255: ConvergenceWarning: Solver terminated early (max_iter=
10000). Consider pre-processing your data with StandardScaler or MinMaxScaler.
  warnings.warn('Solver terminated early (max_iter=%i).')
```

```
In [13]: y_pred = svr.predict(data_X_test)
```

```
In [14]: r2_score(data_y_test, y_pred), mean_absolute_percentage_error(data_y_test, y_pred)
```

```
Out[14]: (0.8007830168563735, 0.15056431505942702)
```

## Деревья решений

```
In [15]: # Обучим дерево и предскажем результаты
from sklearn.tree import DecisionTreeRegressor

tree_regr = DecisionTreeRegressor(random_state=1, max_depth=7).fit(data_X_train, data_y_train)
y_test_predict = tree_regr.predict(data_X_test)
y_test_predict.shape
```

```
Out[15]: (1778,)
```

```
In [16]: r2_score(data_y_test, y_test_predict)
```

```
Out[16]: 0.6626220166750287
```

```
In [17]: list(
    zip(
        data_X_train.columns.values,
        tree_regr.feature_importances_
    )
)
```

```
Out[17]: [('Rooms', 0.005024582426006323),
 ('Type', 0.01650526758281762),
```

```
( 'Distance', 0.16679148277830835),
( 'Postcode', 0.09971740359027761),
( 'Bedroom2', 0.0006393899010610526),
( 'Bathroom', 0.011851644859025727),
( 'Car', 0.0),
( 'Landsize', 0.0506516356530368),
( 'BuildingArea', 0.4239520690926708),
( 'YearBuilt', 0.129300603225875),
( 'Latitude', 0.06167573602026102),
( 'Longitude', 0.032645389732301903),
( 'Propertycount', 0.0012447951383578269)]
```

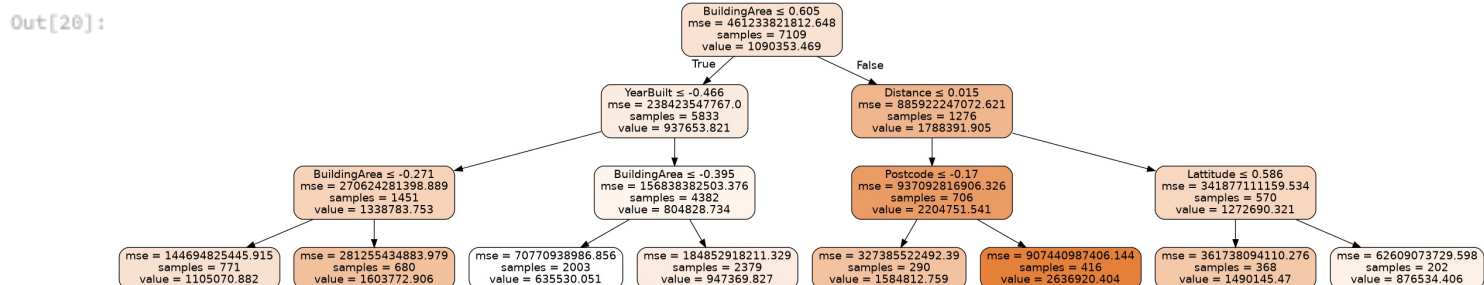
```
In [18]: # Визуализация дерева
from sklearn.tree import export_graphviz
from io import StringIO
import pydot

def get_png_tree(tree_model_param, feature_names_param):
    dot_data = StringIO()
    export_graphviz(
        tree_model_param, out_file=dot_data,
        feature_names=feature_names_param, filled=True,
        rounded=True, special_characters=True
    )
    (graph,) = pydot.graph_from_dot_data(dot_data.getvalue())
    return graph.create_png()
```

```
In [19]: tree_regr_depth3 = DecisionTreeRegressor(
    random_state=1,
    max_depth=3
).fit(data_X_train, data_y_train)
tree_regr_depth3
```

Out[19]: DecisionTreeRegressor(max\_depth=3, random\_state=1)

```
In [20]: from IPython.display import Image
Image(get_png_tree(tree_regr_depth3, [x for x in data.columns if x != "Price"]))
```



In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js